

Link layer, LANs: roadmap

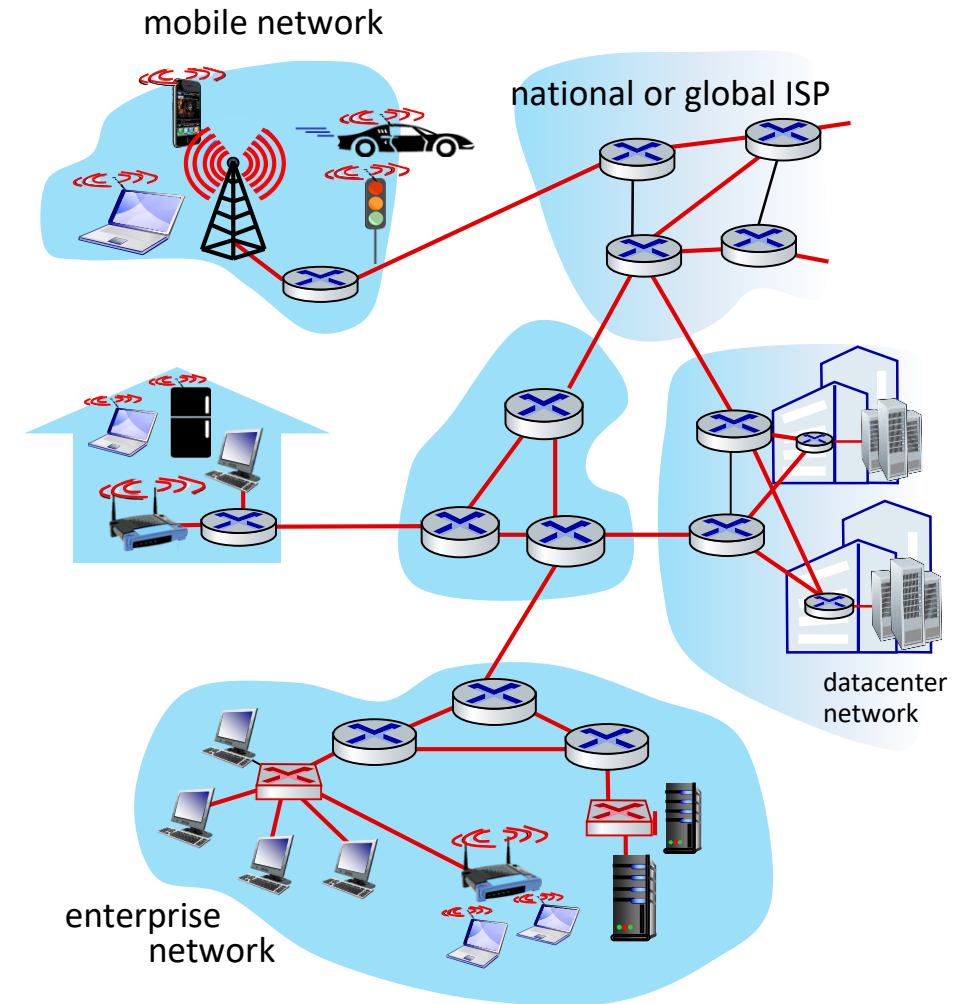
- introduction
- error detection
- LANs
 - addressing, ARP
 - Ethernet
 - switches
 - VLANs

Link layer: introduction

terminology:

- hosts and routers: nodes
- communication channels that connect adjacent nodes along communication path: links
 - wired
 - wireless
 - LANs
- layer-2 packet: *frame*, encapsulates datagram

*link layer has responsibility of transferring datagram from one node to **physically adjacent** node over a link*





Link layer: context

- datagram transferred by different link protocols over different links:
 - e.g., WiFi on first link, Ethernet on next link
- each link protocol provides different services
 - e.g., may or may not provide reliable data transfer over link



Link layer: services

- **framing, link access:**
 - encapsulate datagram into frame, adding header, trailer
 - channel access
 - “Medium Access Control-MAC” addresses in frame headers identify source, destination
- **reliable delivery between adjacent nodes**
 - seldom used on low bit-error links (fiber, coax, and twisted-pair copper links)
 - often used for wireless links: high error rates
 - Q: why both link-level and end-end reliability?
 - correcting an error locally, on the link where the error occurs, rather than forcing an end-to-end retransmission of the data by a transport or application-layer protocol



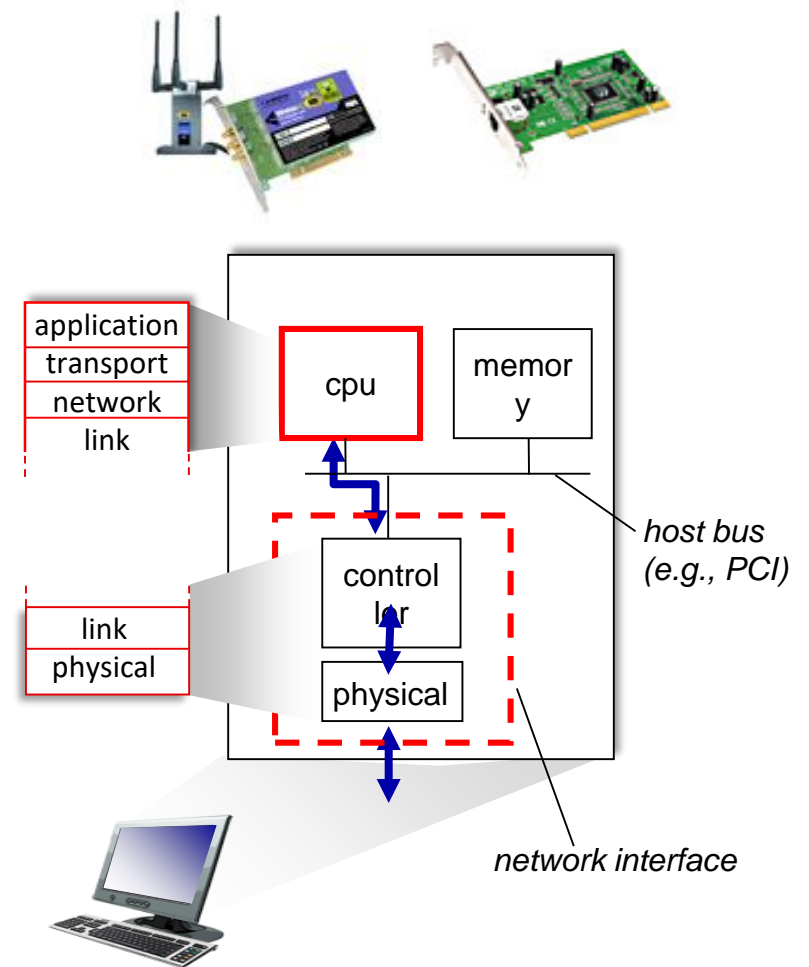
Link layer: services (more)

- **flow control:**
 - pacing between adjacent sending and receiving nodes
- **error detection:**
 - errors caused by signal attenuation, noise (bit errors, 0 instead of 1 or vice versa)
- **error correction:**
 - receiver identifies *and corrects* bit error(s) without retransmission
- **half-duplex and full-duplex:**
 - with half duplex, nodes at both ends of link can transmit, but not at same time




Where is the link layer implemented?

- in each-and-every host
- link layer implemented in *network (adapter) or interface card* (NIC) or on a chip (hardware)
 - Ethernet, WiFi card or chip
- attaches into host's system buses
- combination of hardware, software, firmware



Interfaces communicating



sending side:

- encapsulates datagram in frame
- adds error checking bits, reliable data transfer, flow control, etc.

receiving side:

- looks for errors, reliable data transfer, flow control, etc.
- extracts datagram, passes to upper layer at receiving side

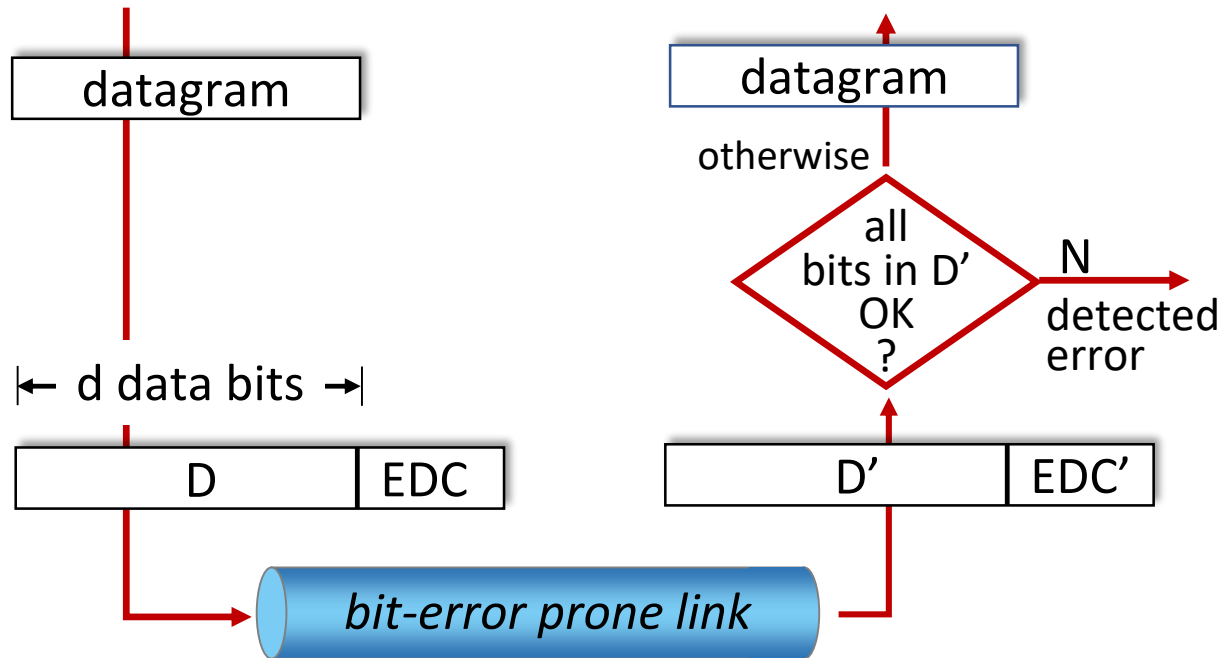
Link layer, LANs: roadmap

- introduction
- **error detection**
- LANs
 - addressing, ARP
 - Ethernet
 - switches
 - VLANs

Error detection

EDC: error detection and correction bits (e.g., redundancy)

D: data protected by error checking, may include header fields



Error detection not 100% reliable!


- protocol may miss some errors, but rarely
- larger EDC field yields better detection and correction

Parity checking

the simplest form of error detection

single bit parity:

- detect single bit errors



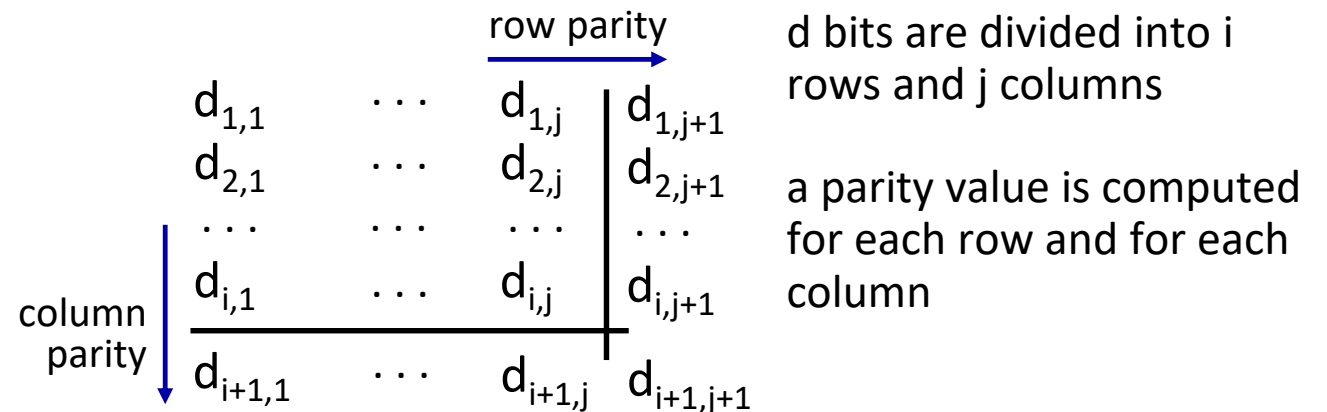
Even parity: set parity bit so there is an even number of 1's (in d+1 bits)

total number of ones is counted, if it is not even, an error is likely to have occurred

only detects a single error, if two bits are flipped, it won't catch it

two-dimensional bit parity:

- detect *and correct* single bit errors



no errors:

1	0	1	0	1	1
1	1	1	1	0	0
0	1	1	1	0	1
0	0	1	0	1	0

detected
and
correctable
single-bit
error:

1	0	1	0	1	1
1	0	1	1	0	0
0	1	1	1	0	1
0	0	1	0	1	0

parity error

parity error

Internet checksum

Goal: detect errors (*i.e.*, flipped bits) in transmitted segment

sender:


- treat contents of UDP segment (including UDP header fields and IP addresses) as sequence of 16-bit integers
- **checksum:** addition (one's complement sum) of segment content
- checksum value put into UDP checksum field

receiver:

- compute checksum of received segment
- check if computed checksum equals checksum field value:
 - not equal - error detected
 - equal - no error detected

Cyclic Redundancy Check (CRC)

- the most powerful method
- **D**: data bits (given, think of these as a binary number)
- **G**: bit pattern (generator), of $r+1$ bits (given)



goal: choose r CRC bits, **R**, such that $\langle D, R \rangle$ exactly divisible by $G \pmod{2}$

- receiver knows G , divides $\langle D, R \rangle$ by G . If non-zero remainder: error detected!
- can detect all burst errors less than $r+1$ bits
- widely used in practice (Ethernet, 802.11 WiFi)

Cyclic Redundancy Check (CRC): example

We want:

$$D \cdot 2^r \text{ XOR } R = nG$$

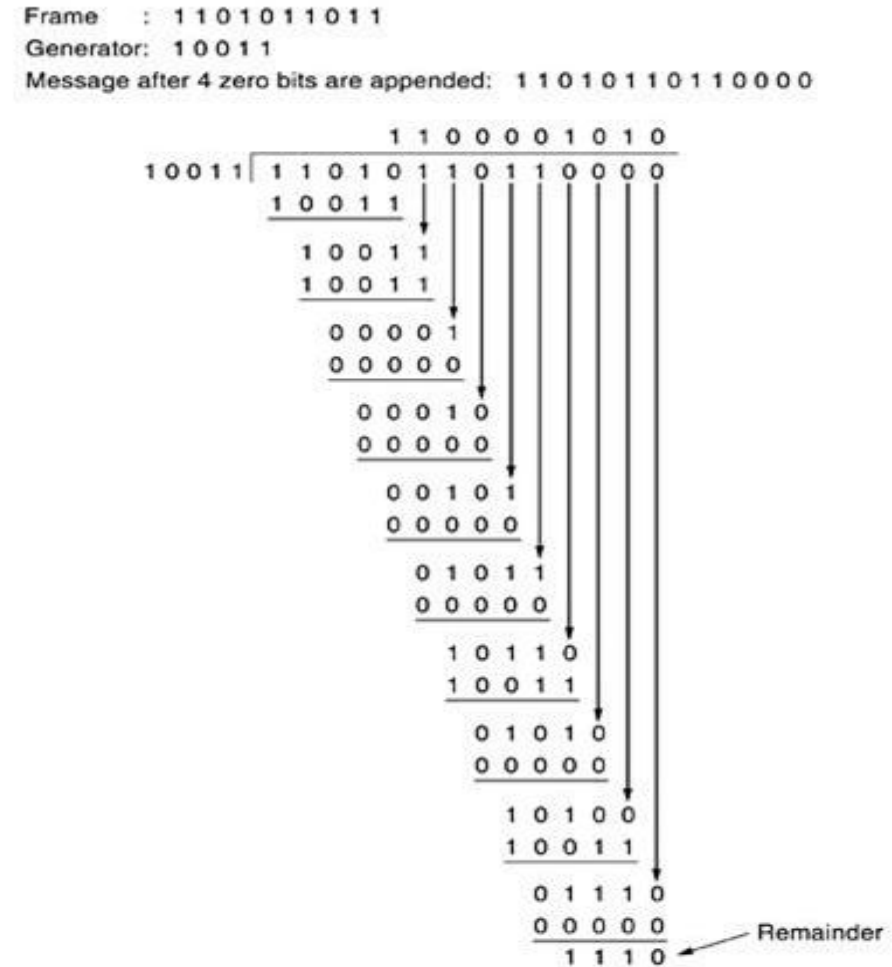
or equivalently:

$$D \cdot 2^r = nG \text{ XOR } R$$

or equivalently:

if we divide $D \cdot 2^r$ by G , want remainder R to satisfy:

$$R = remainder \left[\frac{D \cdot 2^r}{G} \right]$$



Transmitted frame: 1 1 0 1 0 1 1 0 1 1 1 1 1 0

Link layer, LANs: roadmap

- introduction
- error detection, correction
- LANs
 - addressing, ARP
 - Ethernet
 - switches
 - VLANs


MAC addresses

- MAC (or LAN or physical or Ethernet) address:
 - function: used “locally” to get frame from one interface to another physically-connected interface (same subnet, in IP-addressing sense)
 - 48-bit MAC address (for most LANs) burned in NIC ROM, also sometimes software settable
 - e.g.: 1A-2F-BB-76-09-AD
 - hexadecimal (base 16) notation
(each “numeral” represents 4 bits)*

MAC addresses

each interface on LAN

- has unique 48-bit **MAC** address
- has a locally unique 32-bit IP address




MAC addresses

- MAC address allocation administered by IEEE
- manufacturer buys portion of MAC address space (to assure uniqueness)
- MAC flat address: portability
 - can move interface from one LAN to another
 - recall IP address *not* portable: depends on IP subnet to which node is attached

ARP: address resolution protocol

Question: how to determine interface's MAC address, knowing its IP address? (IP addresses are NOT recognized by hardware)

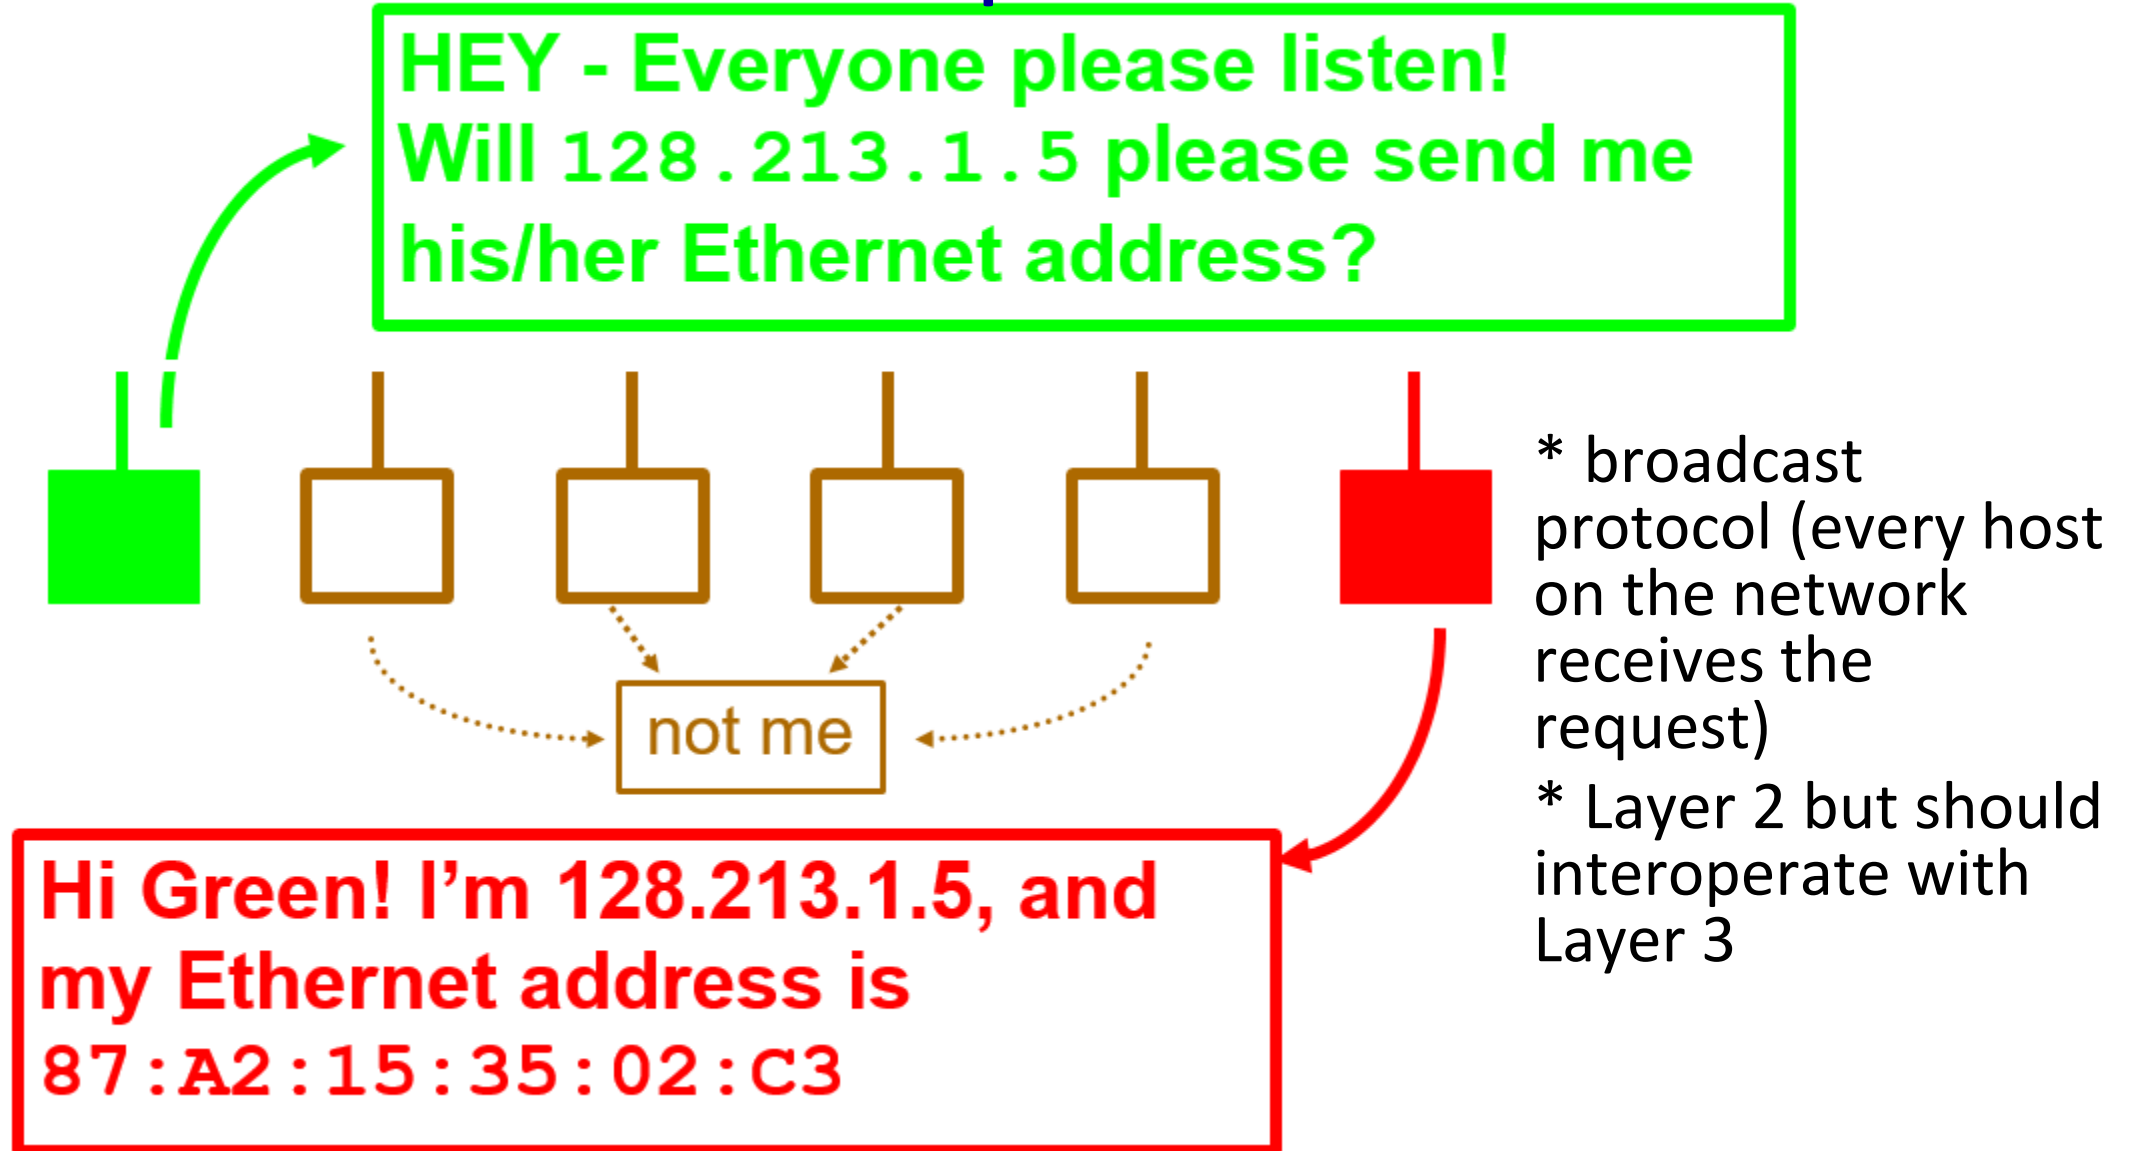
Address Resolution: process of finding the hardware address (physical address) of a host given the IP address (logical address)



ARP table: each IP node (host, router) on LAN has table

- IP/MAC address mappings for some LAN nodes:
< IP address; MAC address; TTL >
- TTL (Time To Live): time after which address mapping will be forgotten

ARP: address resolution protocol



ARP protocol in action

example: A wants to send datagram to B

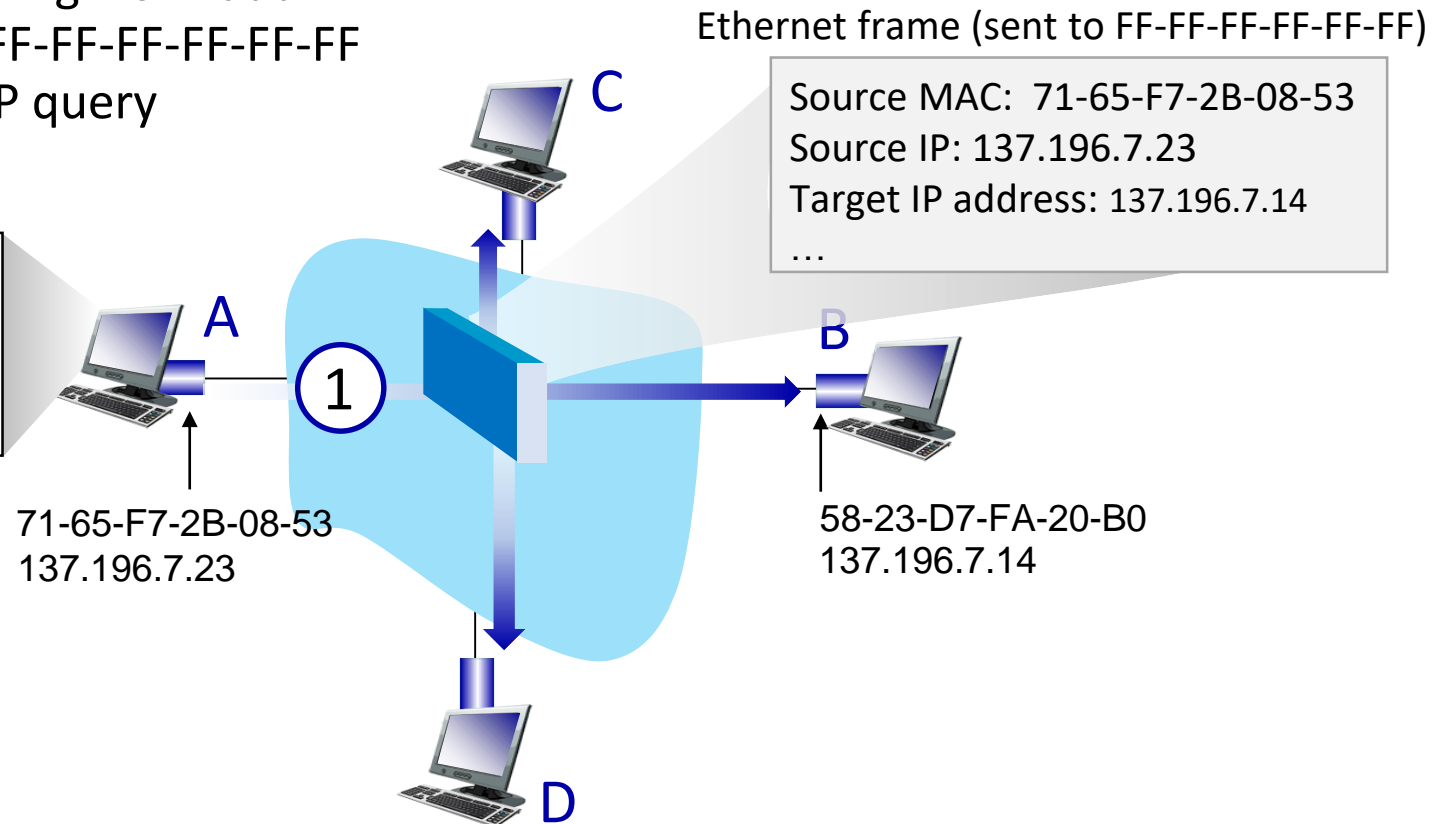
- B's MAC address not in A's ARP table, so A uses ARP to find B's MAC address

A broadcasts ARP query, containing B's IP addr

- ①
- destination MAC address = FF-FF-FF-FF-FF-FF
 - all nodes on LAN receive ARP query

ARP table in A

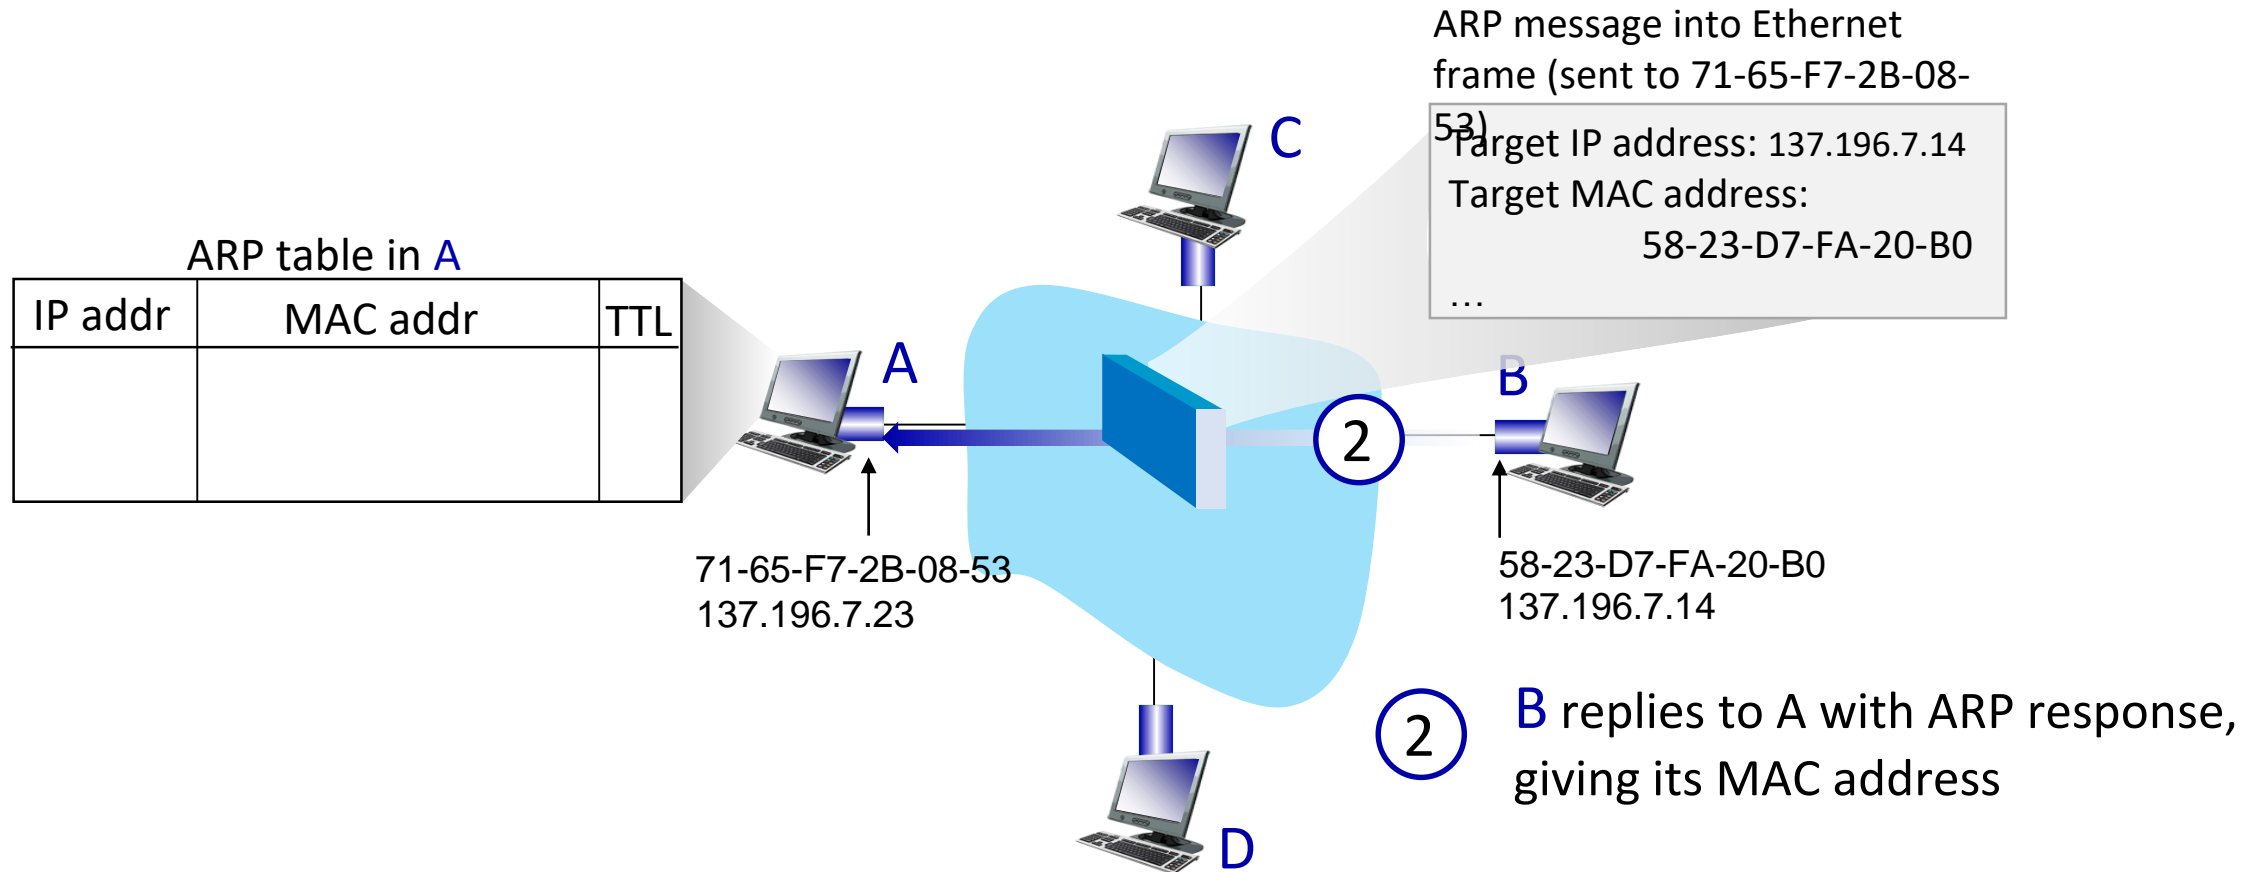
IP addr	MAC addr	TTL



ARP protocol in action

example: A wants to send datagram to B

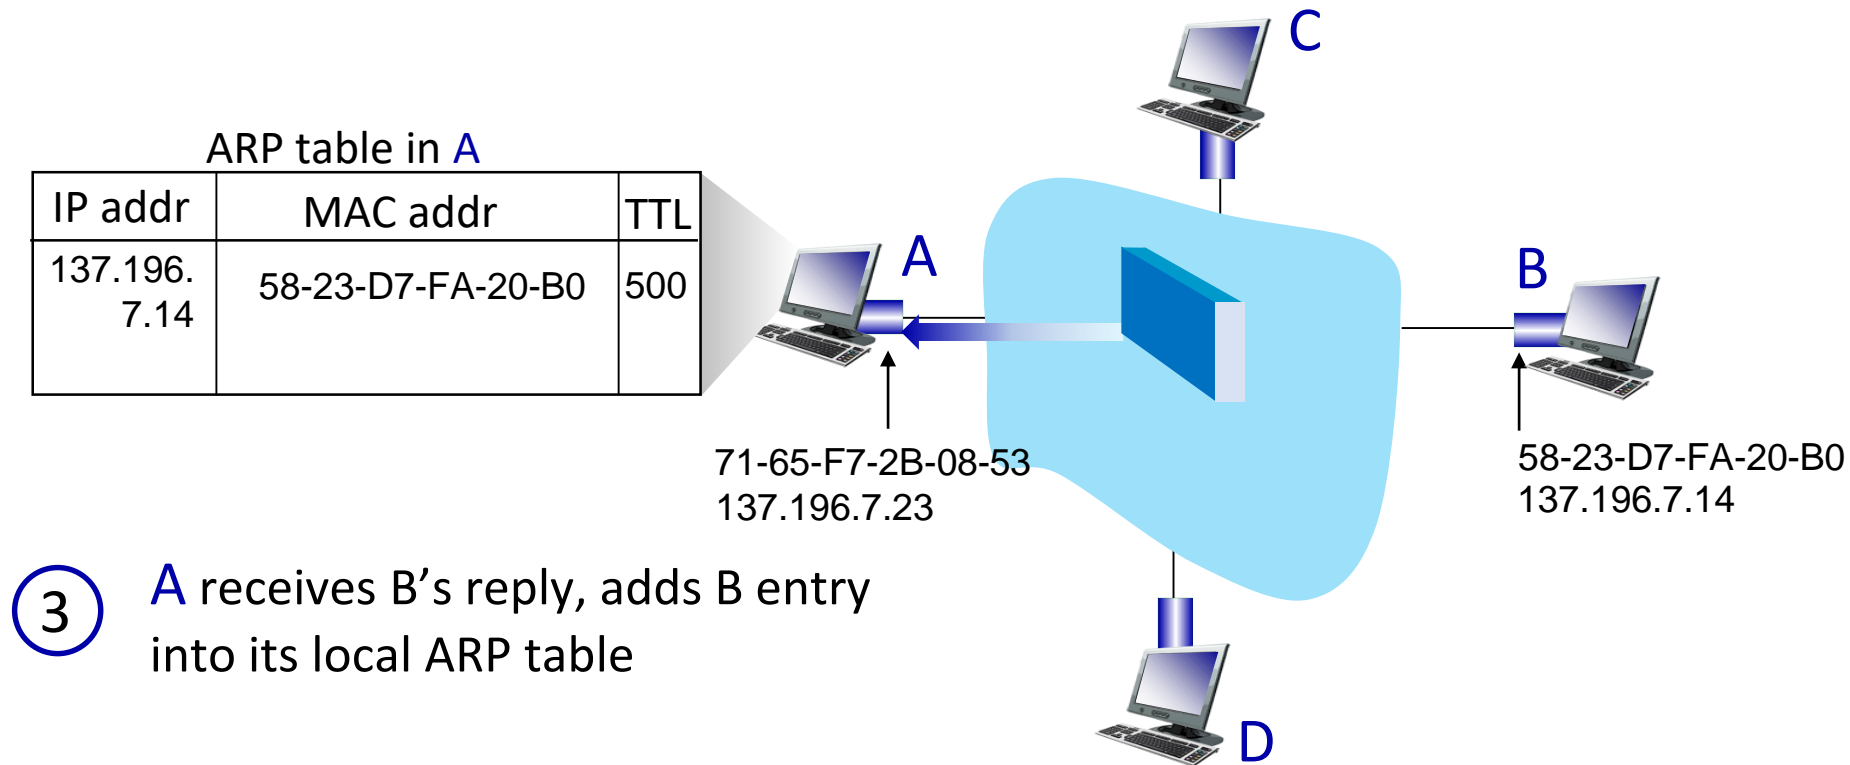
- B's MAC address not in A's ARP table, so A uses ARP to find B's MAC address



ARP protocol in action

example: A wants to send datagram to B


- B's MAC address not in A's ARP table, so A uses ARP to find B's MAC address



Routing to another subnet: addressing


walkthrough: sending a datagram from *A* to *B* via *R*

- focus on addressing – at IP (datagram) and MAC layer (frame) levels
- assume that:
 - A knows B's IP address
 - A knows IP address of first hop router, R
 - A knows R's MAC address



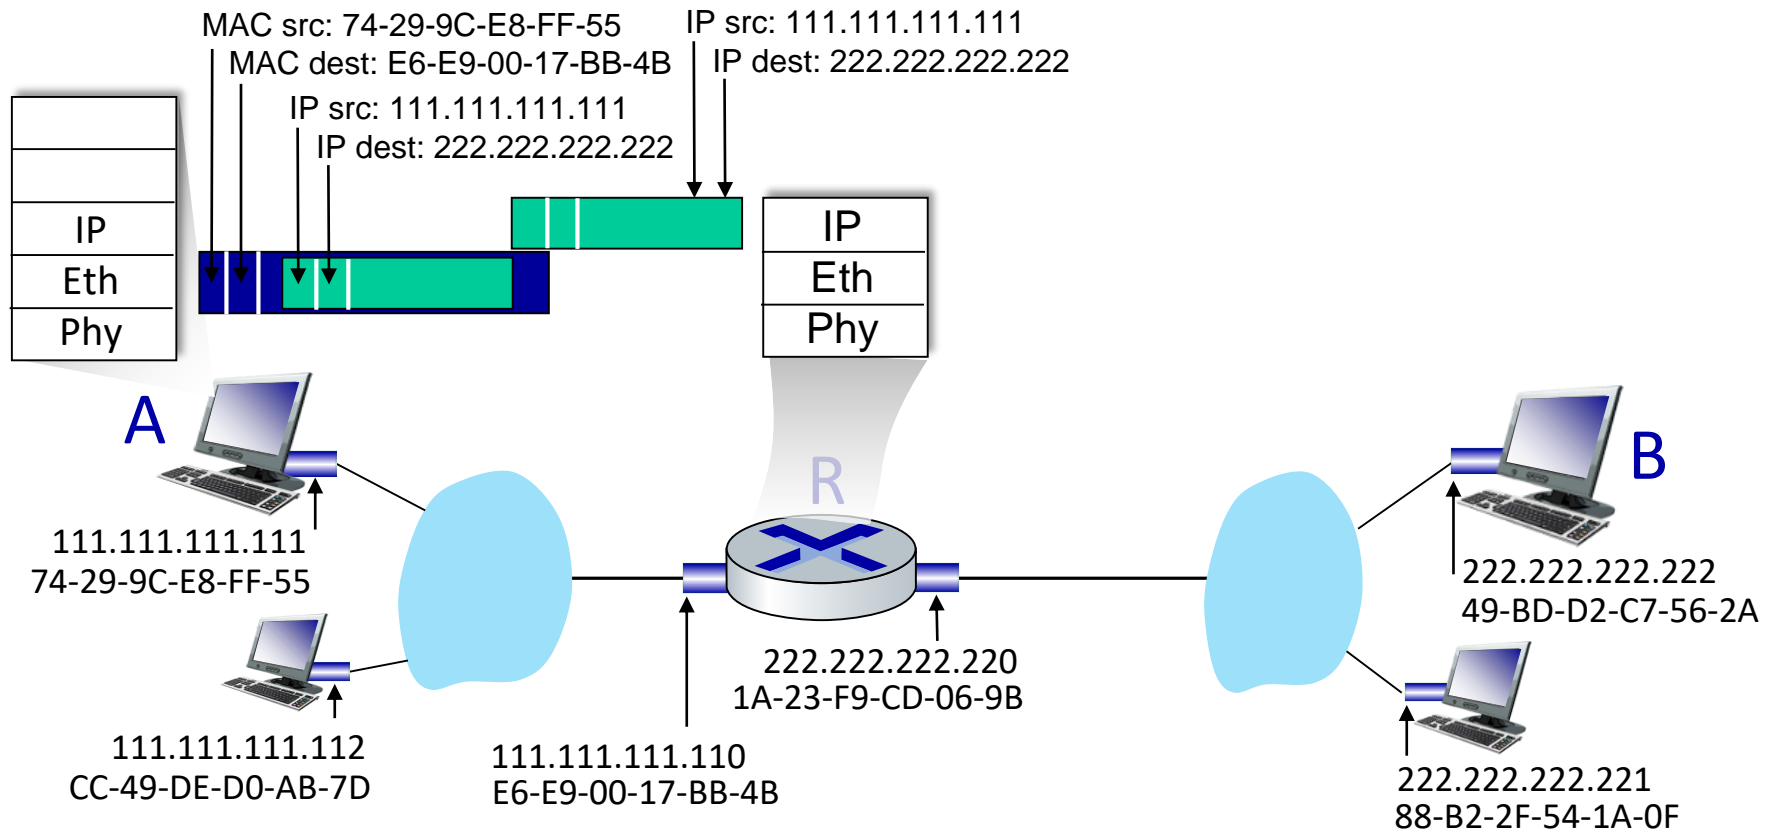
Routing to another subnet: addressing

- A creates IP datagram with IP source A, destination B
- A creates link-layer frame containing A-to-B IP datagram
 - **R's** MAC address is frame's destination




Routing to another subnet: addressing

- frame sent from A to R
- frame received at R, datagram removed, passed up to IP




Routing to another subnet: addressing

- R determines outgoing interface, passes datagram with IP source A, destination B to link layer
- R creates link-layer frame containing A-to-B IP datagram. Frame destination address: B's MAC address




Routing to another subnet: addressing

- R determines outgoing interface, passes datagram with IP source A, destination B to link layer
- R creates link-layer frame containing A-to-B IP datagram. Frame destination address: B's MAC address
- **transmits link-layer frame**



Routing to another subnet: addressing


- B receives frame, extracts IP datagram destination B
- B passes datagram up protocol stack to IP



Routing to another subnet: addressing

PC1 Sends a Packet to PC2

Because PC2 is on different network, I will encapsulate the packet and send it to the router on MY network. Let me find that MAC address....




Layer 2 Data Link Frame

Dest. MAC	Source MAC	Type	Source IP	Dest. IP	IP fields	Data	Trailer
00-10	0A-10	800	192.168.1.10	192.168.4.10			

Packet's Layer 3 data

PC1's ARP Cache for R1	
IP Address	MAC Address
192.168.1.1	00-10

Routing to another subnet: addressing



Layer 2 Data Link Frame

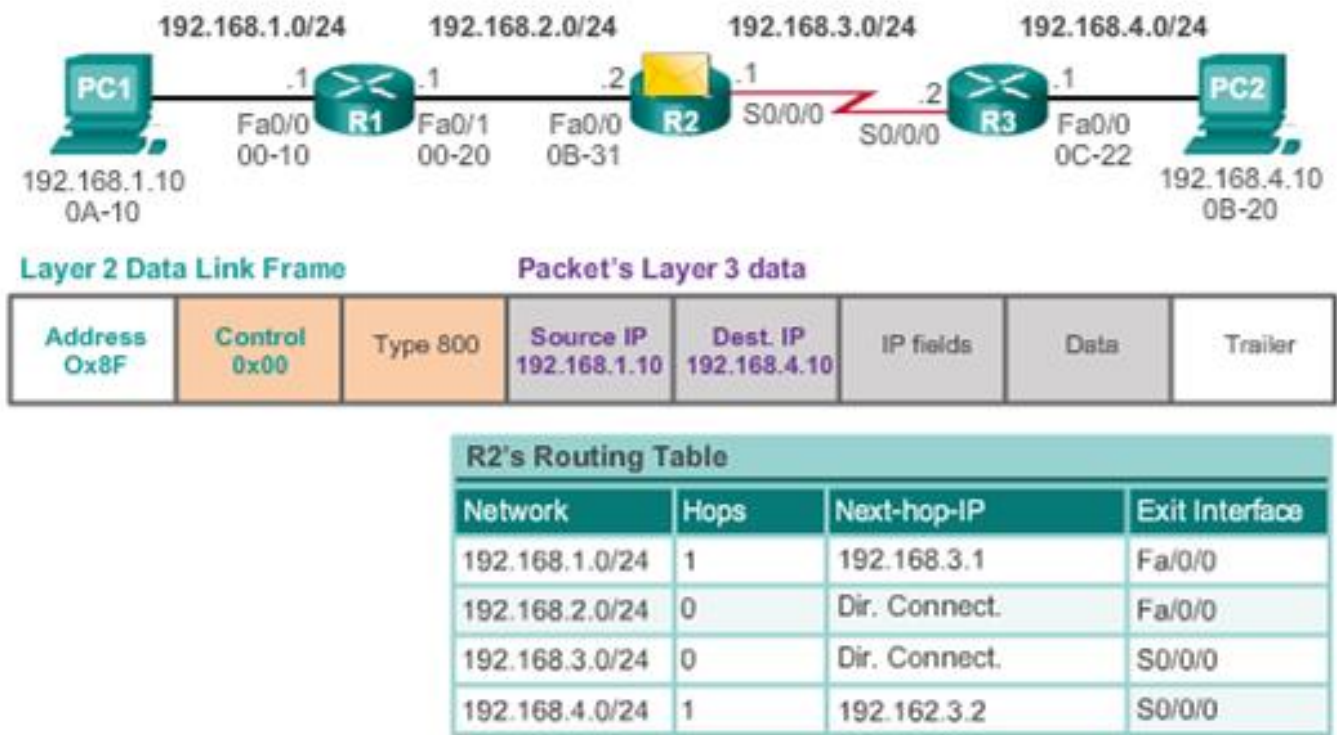
Packet's Layer 3 data

Dest. MAC 0B-31	Source MAC 00-20	Type 800	Source IP 192.168.1.10	Dest. IP 192.168.4.10	IP fields	Data	Trailer
--------------------	---------------------	----------	---------------------------	--------------------------	-----------	------	---------

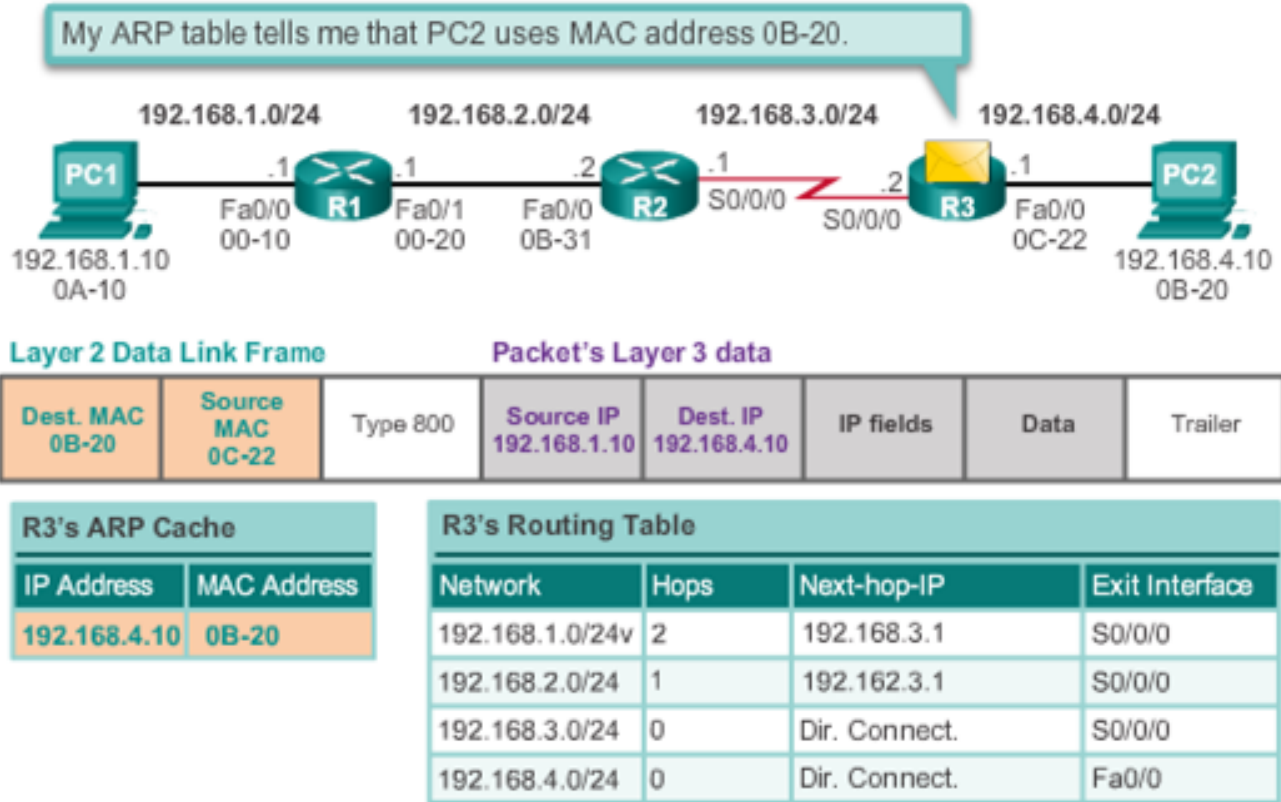
R1's Routing Table

Network	Hops	Next-hop-IP	Exit Interface
192.168.1.0/24	0	Dir. Connect.	Fa0/0
192.168.2.0/24	0	Dir. Connect.	Fa0/1
192.168.3.0/24	1	192.168.2.2	Fa0/1
192.168.4.0/24	2	192.168.2.2	Fa0/1

Routing to another subnet: addressing



Routing to another subnet: addressing



Link layer, LANs: roadmap

- introduction
- error detection, correction
- **LANs**
 - addressing, ARP
 - **Ethernet**
 - switches
 - VLANs

Ethernet

“dominant” wired LAN technology:

- first widely used LAN technology
- simpler, cheap

Ethernet frame structure


sending interface encapsulates IP datagram (or other network layer protocol packet) in **Ethernet frame**



preamble:

- used to synchronize receiver, sender clock rates (10 Mbps, 100 Mbps, or 1 Gbps)
- 8 bytes

Ethernet frame structure (more)



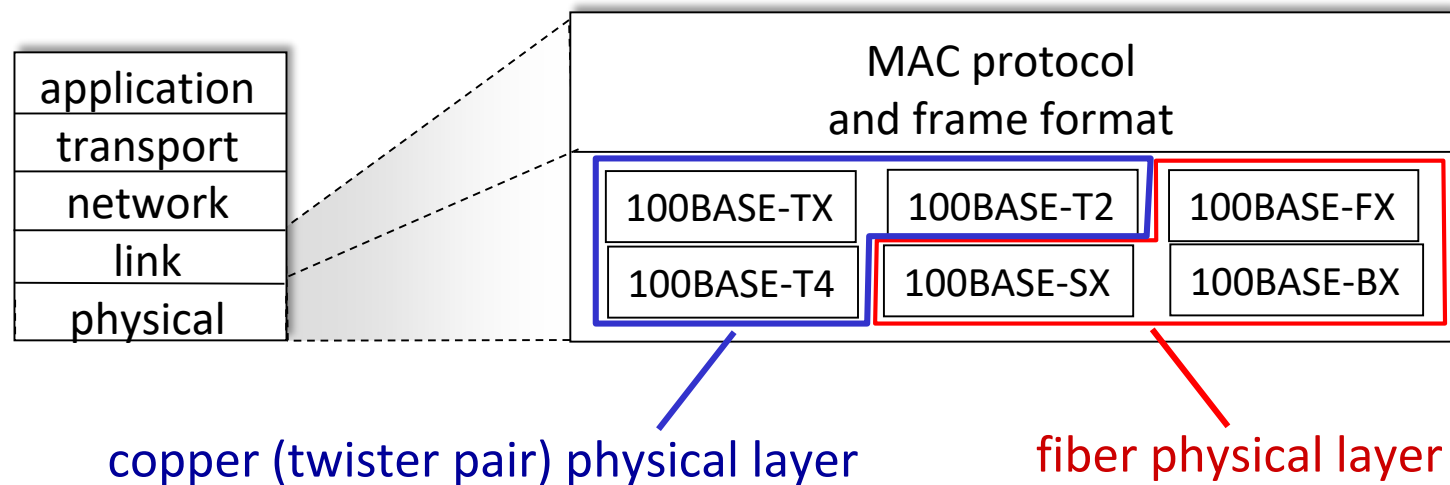
- **addresses:** 6 byte source, destination MAC addresses
 - if adapter receives frame with matching destination address, or with broadcast address (e.g., ARP packet), it passes data in frame to network layer protocol
 - otherwise, adapter discards frame
- **type:** indicates higher layer protocol
 - mostly IP but others possible, e.g., Novell IPX, AppleTalk
 - used to demultiplex up at receiver
- **CRC:** cyclic redundancy check at receiver
 - error detected: frame is dropped

Ethernet: unreliable, connectionless

- **connectionless:** no handshaking between sending and receiving NICs
- **unreliable:** receiving NIC doesn't send ACKs
 - data in dropped frames recovered only if initial sender uses higher layer reliable data transfer (e.g., TCP), otherwise dropped data lost

802.3 Ethernet standards: link & physical layers

- *many* different Ethernet standards
 - common MAC protocol and frame format
 - different speeds: 2 Mbps, 10 Mbps, 100 Mbps, 1Gbps, 10 Gbps, 40 Gbps
 - different physical layer media: fiber, cable



Link layer, LANs: roadmap


- introduction
- error detection, correction
- LANs
 - addressing, ARP
 - Ethernet
 - switches
 - VLANs

Ethernet switch

- Switch is a **link-layer** device: takes an *active* role
 - store, forward Ethernet frames
 - examine incoming frame's MAC address, *selectively* forward frame to one-or-more outgoing links when frame is to be forwarded on segment
- **transparent**: hosts *unaware* of presence of switches
- **plug-and-play, self-learning**

Switch: multiple simultaneous transmissions

- hosts have dedicated, direct connection to switch
- switches buffer packets
- Ethernet protocol used on *each* incoming link, so:
 - no collisions; full duplex
- **switching:** A-to-A' and B-to-B' can transmit simultaneously, without collisions (but A-to-A' and C to A' can not happen simultaneously)




switch with six
interfaces (1,2,3,4,5,6)

Switch forwarding table

Q: how does switch know A' reachable via interface 4, B' reachable via interface 5?

A: each switch has a **switch table**, each entry:


- (MAC address of host, interface to reach host)
- like a routing table



Switch: self-learning

how are entries created, maintained in switch table?

- switch *learns* which hosts can be reached through which interfaces
- when frame received, switch “learns” location of sender: incoming LAN segment
- records sender/location pair in switch table




Switch: frame filtering/forwarding

when frame received at switch:

1. record incoming link, MAC address of sending host
2. index switch table using MAC destination address
3. if entry found for destination
 then {
 if destination on segment from which frame arrived
 then drop frame
 else forward frame on interface indicated by entry
 }
 else flood /* forward on all interfaces except arriving interface */

Interconnecting switches


self-learning switches can be connected together:



Q: sending from A to G - how does S_1 know to forward frame destined to G via S_4 and S_3 ?

- A: self learning (works exactly the same as in single-switch case)

Small institutional network



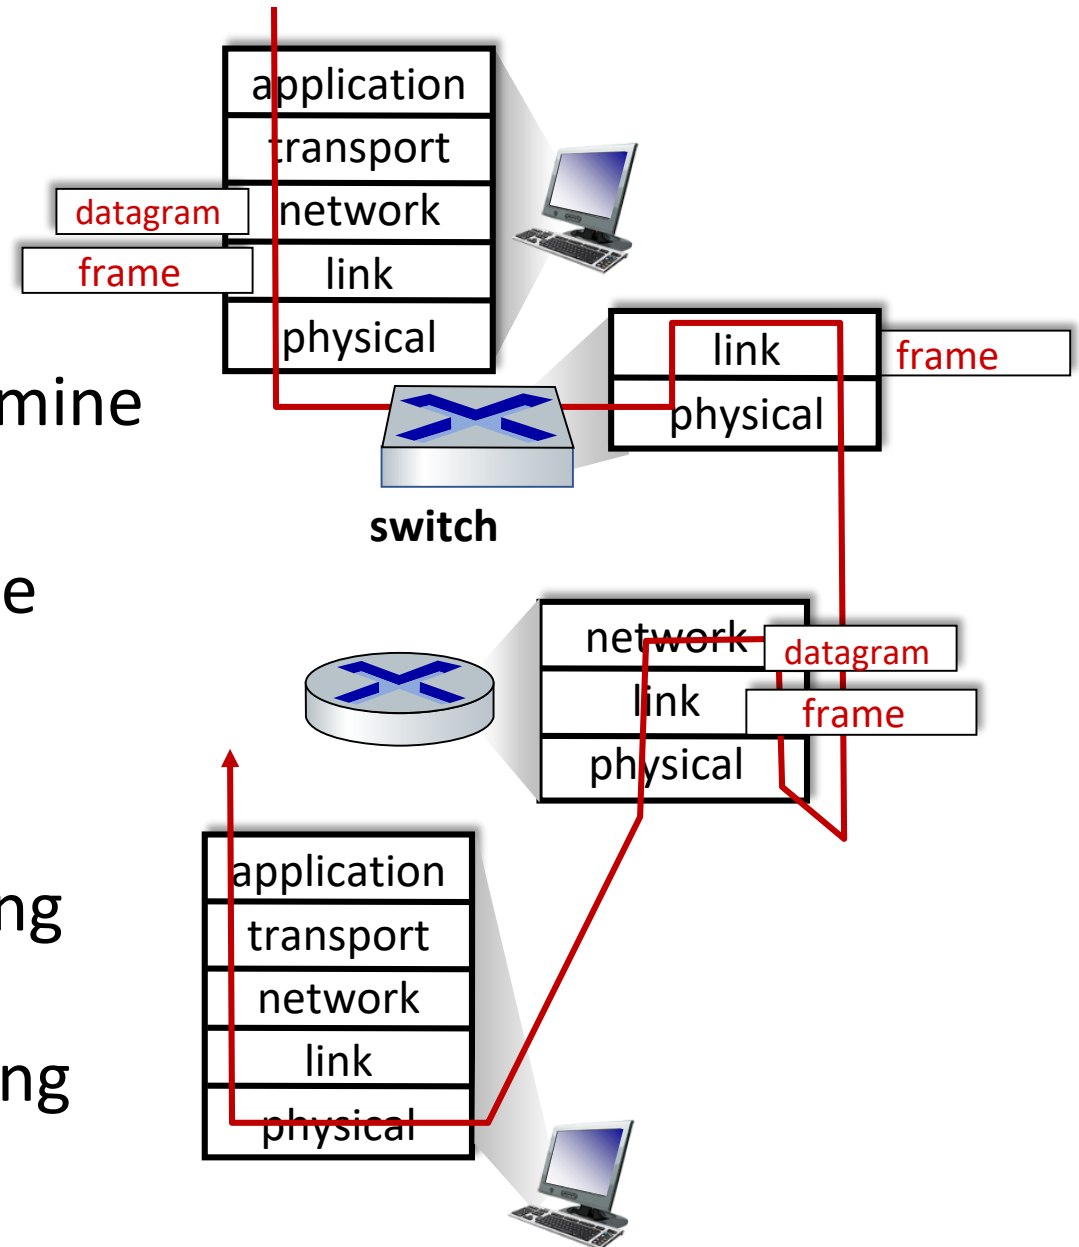
Switches vs. routers

both are store-and-forward:

- *routers*: network-layer devices (examine network-layer headers)
- *switches*: link-layer devices (examine link-layer headers)

both have forwarding tables:

- *routers*: compute tables using routing algorithms, IP addresses
- *switches*: learn forwarding table using flooding, learning, MAC addresses



Link layer, LANs: roadmap

- introduction
- error detection, correction
- **LANs**
 - addressing, ARP
 - Ethernet
 - switches
 - **VLANs**

Virtual LANs (VLANs): motivation

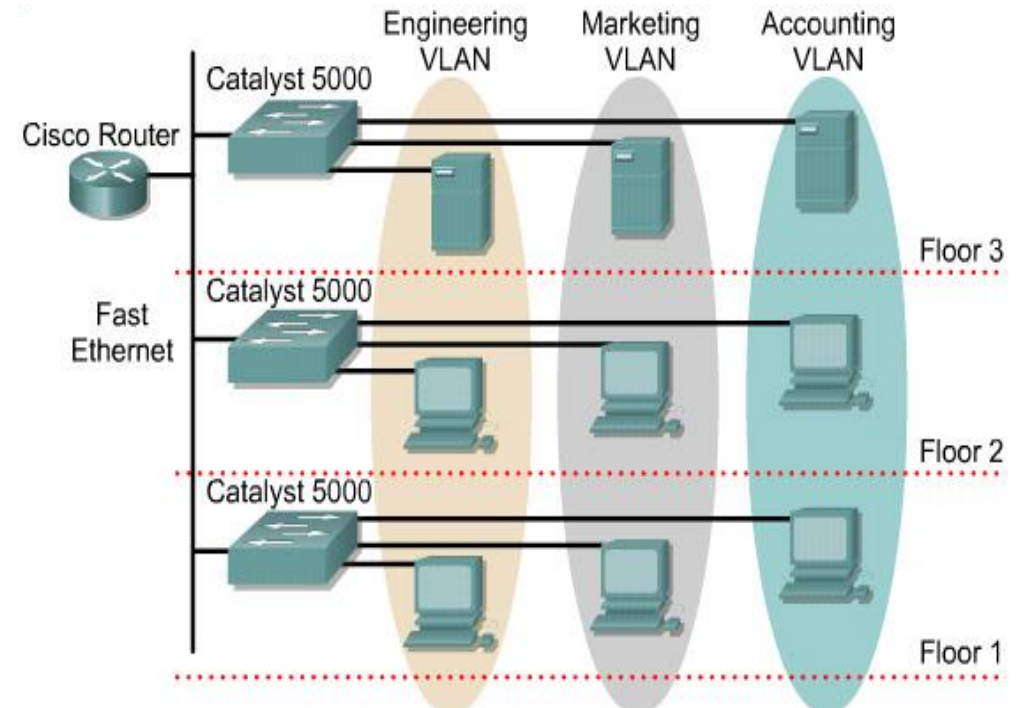

Q: what happens as LAN sizes scale, users change point of attachment?

VLANs logically segment switched networks

Packets are only switched between ports that are designated for the same VLAN

The key benefit: organizing the LAN logically instead of physically

Layer 3 routing allows the router to send packets to the three different VLANs (interconnect)




Port-based VLANs


Virtual Local Area Network (VLAN)

switch(es) supporting VLAN capabilities can be configured to define multiple *virtual* LANS over single physical LAN infrastructure.

port-based VLAN: switch ports grouped (by switch management software) so that *single* physical switch



... operates as **multiple** virtual switches



Port-based VLANs

- **traffic isolation:** frames to/from ports 1-8 can *only* reach ports 1-8
 - can also define VLAN based on MAC addresses of endpoints, rather than switch port
- **dynamic membership:** ports can be dynamically assigned among VLANs
- **forwarding between VLANs:** done via routing (just as with separate switches)
 - in practice vendors sell combined switches plus routers

